


Indexing

Ramesh S



Index types

- Default _id
- Single
- Compound
- Multi key
- Geospatial
- Text
- Hashed



Let's create a **biggg** collection

use big

```
for(i=0;i<1000000;i++){print(i);db.big.insert({a:i,b:i+2,c:i+4})}
```

```
db.big.count()
```

Let's check its size

`db.big.dataSize()`

`db.big.storageSize()`

`db.big.getIndexes()`

Let's create some indexes


```
db.big.createIndex({a:1})
```

```
db.big.createIndex({a:1,b:1},{background:true})
```

```
db.big.createIndex({a:1,b:1,c:1},{background:true})
```

```
db.big.createIndex({c:1},{unique:true})
```

```
db.big.getIndexes()
```



Let's check the size of indexes

`db.big.totalIndexSize()`


`db.big.dataSize()`

`db.big.totalSize()`

`db.big.storageSize()`

Checking how good an index is

```
db.big.find({a:89}).explain("executionStats")
```



Let's Intersect them

```
db.big.find( {a: { $gt: 9 } , c: 14} )
```

Returns a single document having a=12, b=12 and c=14

Working with indexes

`db.big.getIndexes()`

`db.big.dropIndexes()`

`db.big.reIndex()`

`db.big.dropIndex(index)`

TTL indexes

```
db.events.createIndex( { "loggedAt": 1 }, { expireAfterSeconds: 864000 } )
```

```
db.events.insert({ "loggedAt": new Date(), "mesg": "Disk full" })
```

```
db.events.insert({ "loggedAt": new Date(), "mesg": "Server not responding" })
```

TTL indexes

```
db.events.createIndex( { "eraseAt": 1 }, { expireAfterSeconds: 0 } )
```

```
db.events.insert({ "eraseAt": new Date('June 30, 2016 14:00:00'), "mesg": "Disk full" })
```

```
db.events.insert({ "eraseAt": new Date('June 30, 2016 14:00:00'),  
                  "mesg": "Server not responding" })
```


Index usage stats

```
db.big.aggregate( { $indexStats: {} } )
```

GeoJSON Objects

To specify GeoJSON data, use an embedded document with:

- a field named `type` that specifies the GeoJSON object type and
- a field named `coordinates` that specifies the object's coordinates.
- If specifying latitude and longitude coordinates, list the longitude first and then latitude:
 - Valid longitude values are between -180 and 180, both inclusive.
 - Valid latitude values are between -90 and 90, both inclusive.




GeoJSON Objects

```
location: {  
  type: "Point",  
  coordinates: [-73.856077, 40.848447]  
}
```

Legacy Coordinate Pairs

- To calculate distances on a Euclidean plane, store your location data as legacy coordinate pairs and use a 2d index.
- MongoDB supports spherical surface calculations on legacy coordinate pairs via a 2dsphere index by converting the data to the GeoJSON Point type.



Legacy Coordinate Pairs

location: [-73.856077, 40.848447]

- Valid longitude values are between -180 and 180, both inclusive.
- Valid latitude values are between -90 and 90, both inclusive.

Geospatial Indexes

MongoDB provides the following geospatial index types to support the geospatial queries.

- 2dsphere

```
db.big.createIndex( { <location field> : "2dsphere" } )
```

where the <location field> is a field whose value is either a GeoJSON object or a legacy coordinates pair.

- 2d

```
db.collection.createIndex( { <location field> : "2d" } )
```

where the <location field> is a field whose value is a legacy coordinates pair.

Hashed Indexes

- Hashed indexes maintain entries with hashes of the values of the indexed field.
- Hashed indexes support sharding using hashed shard keys.
- Hashed based sharding uses a hashed index of a field as the shard key to partition data across your sharded cluster.
- Using a hashed shard key to shard a collection results in a more random distribution of data.

Hashed Indexes

To create a hashed index, specify hashed as the value of the index key

```
db.big.createIndex( { _id: "hashed" } )
```

Hashed Indexes

- MongoDB supports hashed indexes of any single field.
- The hashing function collapses embedded documents and computes the hash for the entire value, but does not support multi-key (i.e. arrays) indexes.
- You may not create compound indexes that have hashed index fields or specify a unique constraint on a hashed index;
- You can create both a hashed index and an ascending/descending (i.e. non-hashed) index on the same field: MongoDB will use the scalar index for range queries.

List all Hashed Indexes

```
db.adminCommand("listDatabases").databases.forEach(function(d){  
  let mdb = db.getSiblingDB(d.name);  
  mdb.getCollectionInfos({ type: "collection" }).forEach(function(c){  
    let currentCollection = mdb.getCollection(c.name);  
    currentCollection.getIndexes().forEach(function(idx){  
      let idxValues = Object.values(Object.assign({}, idx.key));  
  
      if (idxValues.includes("hashed")) {  
        print("Hashed index: " + idx.name + " on " + idx.ns);  
        printjson(idx);  
      }  
    });  
  });  
});
```

