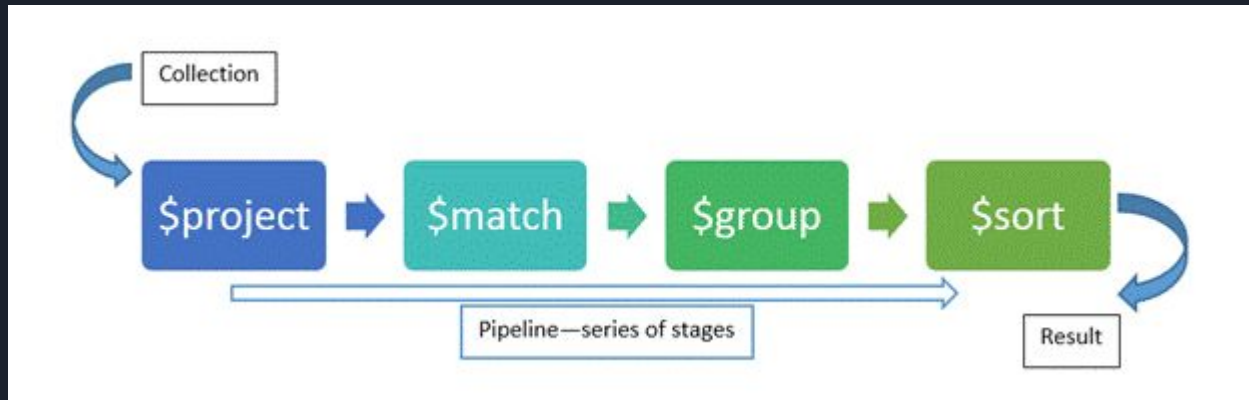




# MongoDB Aggregation

Ramesh S

# Aggregation Pipeline





# Simple Group by

```
db.products.aggregate([
  {$group:
    {
      "_id":"$category",
      "num_of_products":{$sum:1}
    }
  }
])
```





# The Aggregation Pipeline

<code>\$project</code>	select, reshape	1:1
<code>\$match</code>	filter	n:1
<code>\$group</code>	aggregate	n:1
<code>\$sort</code>	sort	1:1
<code>\$skip</code>	skip	n:1
<code>\$limit</code>	limit	n:1
<code>\$unwind</code>		1:n



# Compound Aggregation

```
db.products.aggregate([
  {$group:
    {
      _id:{"maker"  : "$manufacturer",
          "cat"      : "$category"},
      "num_of_products":{$sum:1}
    }
  ])
```





# Using a Document for \_id

- \_id field of a document need not be a scalar value.
- It can be a json document by itself

Example:

```
db.xyz.insert({_id:{name:'ramesh',gender:'m'},profession:"Trainer"})
```



# Aggregation Expressions

`$sum` - used for sum & count

`$avg`

`$min`

`$max`

`$addToSet` - used along with arrays

`$push` - used along with arrays

`$first` - used along with sort

`$last` - used along with sort



# Using \$sum

- To get sum:

```
db.zips.aggregate([  
  $group:  
  { _id:"$state",  
    "population":{"$sum":"$pop"}  
  }  
])
```

- To get count use {\$sum:1}





# Using \$avg

- To get average

```
db.zips.aggregate([  
  $group:  
  { _id:"$state",  
    "average_pop":{"$avg":"$pop"}  
  }  
])
```





# Using \$addToSet

- No parallel in sql world
- Creates an array of values based on the aggregation key.

```
db.zips.aggregate([  
  $group:  
    { _id:"$city","postal_codes":  
      {$addToSet:"$_id"}  
    }  
])
```

- Are both the `_id` fields same?



# Using \$push

- \$push works like \$addToSet
- It does not check for duplicates

```
db.zips.aggregate([  
  $group:  
    { _id:"$city","postal_codes":  
      {$push:"$_id"}  
    }  
])
```





# Using \$max

- Helps you find the maximum value

```
db.zips.aggregate([  
  $group:  
  {  
    "_id": "$state",  
    "pop": { $max: "$pop" }  
  }  
])
```



# Using \$min

- Helps you find the minimum value

```
db.zips.aggregate([  
  $group:  
  {  
    "_id": "$state",  
    "pop": { $min: "$pop" }  
  }  
])
```



# Double Grouping

- Unlike in the sql world you can do double grouping

```
db.class.aggregate([  
  $group:{ _id:{a:"$a", b:"$b"},  
           c:{$max:"$c"}}},  
  {$group:{ _id:"$_id.a",  
           c:{$min:"$c"}}}  
])
```



# Using \$project

- \$project is used to reshape the output of an aggregation
  - remove keys
  - add keys
  - reshape keys
  - use some simple functions on keys like
    - \$toLower
    - \$toUpper
    - \$multiply
    - \$add

# Using \$project

- Example

```
db.products.aggregate([  
  {
```

```
    $project:{
```

```
      _id:0,
```

```
      'maker':{$toLower:"$manufacturer"},
```

```
      'details':{'category':"category",
```

```
        'price':{'$multiply':["$price",10]},
```

```
        'item':'$name'}
```

```
    }
```

```
  ]])
```

create a new  
document  
'details'

omit '\_id' field

rename '\$manufacturer'  
to maker and convert to  
lowercase

multiply '\$price' with 10

rename '\$name' to item







# Using \$match

- \$match works pretty much like find

```
db.zips.aggregate([  
  $match:  
  {  
    pop:{$gt:200000}  
  }  
])
```



# Using \$sort

- Can be used before or after the \$group
- Can be used multiple times
- Can be a memory hog

```
db.zips.aggregate([  
    {$sort:{state:-1}}  
])
```

```
db.zips.aggregate([  
    {$sort:{state:1,city:1}}  
])
```





# \$skip and \$limit

- work exactly the same way as they work with `find`
- Almost always used along with `$sort`
- `$skip` comes first then comes `$limit`.

```
db.zips.aggregate([  
    {$sort:{state:1,city:1}},  
    {$skip:10},  
    {$limit:5}  
])
```



# Example of a Pipelined Aggregate

```
db.zips.aggregate([
  {$match:
    {
      state:"NY"
    }
  },
  {$group:
    {
      _id: "$city",
      population: {$sum:"$pop"},
    }
  },
  {$project:
    {
      _id: 0,
      city: "$_id",
      population: 1,
    }
  },
  {$sort:
    {
      population:-1
    }
  },
  {$skip: 10},
  {$limit: 5}
])
```

# \$first and \$last

- Used to pick up the first or the last of the grouped values

```
db.class.aggregate([  
    {$match:{a:0}},  
    {$sort:{c:-1}},  
    {$group:{_id:"$a",c:{$first:"$c"}}}  
])
```

pick only the first document  
under the grouping





# \$unwind

- Think of it as an opposite to \$push
- Used to process elements in the array, by creating 1 to many documents
- Example:

- Document before unwind

```
{animal:"Cow",eats:["grass","leaves","bananas"]}
```

- Document after unwind

```
{animal:"Cow",eats:"grass"}
```

```
{animal:"Cow",eats:"leaves"}
```

```
{animal:"Cow",eats:"bananas"}
```

- Beware of document explosion

# SQL to Aggregation Mapping Chart

## SQL Terms, Functions, and Concepts

## MongoDB Aggregation Operators

WHERE

\$match

GROUP BY

\$group

HAVING

\$match

SELECT

\$project

ORDER BY

\$sort

LIMIT

\$limit

SUM()

\$sum

COUNT()

\$sum

join

No direct corresponding operator; however, the \$unwind operator allows for somewhat similar functionality, but with fields embedded within the document.





# SQL vs MongoDB Aggregation

- The SQL examples assume two tables, `orders` and `order_lineitem` that join by the `order_lineitem.order_id` and the `orders.id` columns.
- The MongoDB examples assume one collection `orders` that contain documents of the following prototype:

```
{
  cust_id: "abc123",
  ord_date: ISODate("2012-11-02T17:04:11.102Z"),
  status: 'A',
  price: 50,
  items: [ { sku: "xxx", qty: 25, price: 1 },
            { sku: "yyy", qty: 25, price: 1 } ]
}
```







# \$lookup

- Introduced in version 3.2
- Performs a left outer join to an unsharded collection in the *same* database to filter in documents from the “joined” collection for processing.
- To each input document, the \$lookup stage adds a new array field whose elements are the matching documents from the “joined” collection.
- The \$lookup stage passes these reshaped documents to the next stage.





## \$lookup - Single Equality Join

```
db.orders.aggregate([
  {
    $lookup:
    {
      from: "inventory",
      localField: "item",
      foreignField: "sku",
      as: "inventory_docs"
    }
  }
])
```

# SQL vs MongoDB Aggregation

- The MongoDB statements prefix the names of the fields from the documents in the collection orders with a \$ character when they appear as operands to the aggregation operations.

SQL Example	MongoDB Example	Description
<code>SELECT COUNT(*) AS count FROM orders</code>	<pre>db.orders.aggregate( [   { \$group: { _id: null,               count: {                 \$sum: 1 } } } ] )</pre>	Count all records from orders
<code>SELECT SUM(price) AS total FROM orders</code>	<pre>db.orders.aggregate( [   { \$group: { _id: null,               total: {                 \$sum: "\$price" } } } ] )</pre>	Sum the price field from orders

# SQL vs MongoDB Aggregation

SQL Example	MongoDB Example	Description
<pre>SELECT cust_id,        SUM(price) AS total FROM orders GROUP BY cust_id</pre>	<pre>db.orders.aggregate( [   { \$group: { _id:     "\$cust_id",               total: { \$sum: "\$price" } } } ] )</pre>	For each unique cust_id, sum the price field.
<pre>SELECT cust_id,        SUM(price) AS total FROM orders GROUP BY cust_id ORDER BY total</pre>	<pre>db.orders.aggregate( [   { \$group: { _id:     "\$cust_id",               total: { \$sum: "\$price" } } },   { \$sort: { total: 1 } } ] )</pre>	For each unique cust_id, sum the price field, results sorted by sum.
<pre>SELECT cust_id,        ord_date,        SUM(price) AS total FROM orders GROUP BY cust_id, ord_date</pre>	<pre>db.orders.aggregate( [   { \$group: { _id: {     cust_id: "\$cust_id",     ord_date: "\$ord_date" },               total: { \$sum: "\$price" } } } ] )</pre>	For each unique cust_id, ord_date grouping, sum the price field.





# Limitations of MongoDB Aggregations

- Limited to 16MB
- Cannot use more than 10% of the memory on the host machine.